

Local Coding-Model Configurations on a Single DGX Spark:

A Systems and Methodology Case Study

Mani Malarvannan

mani.malarvannan@cdw.com

<https://github.com/mani-mal/spark-coder-bench>

Abstract

We report a reproducible case study of serving and evaluating three open-weight Mixture-of-Experts (MoE) coding models locally on a single NVIDIA DGX Spark (GB10, consumer Grace Blackwell, `sm_121a`, 128 GB unified memory): `gpt-oss-120b` [10] (MXFP4, vLLM), `Qwen3-Coder-30B-A3B` [11] (bf16, vLLM), and NVIDIA Nemotron-3-Super-120B-A12B [12] (NVFP4, TensorRT-LLM). This is deliberately *not* a model-ranking paper. On this hardware each model couples to exactly one serving runtime, one quantization, and one agent scaffold, so every score describes a *deployment configuration*, not intrinsic model quality. Our contributions are three: (1) a serving-feasibility matrix for GB10—the headline result is negative and structural: no model in our set serves on both available runtimes, which makes a matched same-model cross-runtime comparison infeasible on this box; (2) a three-layer local evaluation harness (SWE-bench Verified agentic repair on a disclosed ARM64-buildable subset, a from-scratch full-stack app build scored by an automated API contract, and LiveCodeBench single-shot generation) with full raw-artifact retention and deterministic rescoring; and (3) a set of measurement lessons drawn from concrete failures. The sharpest lesson is a single harness-validity error that compressed a real $\sim 4\times$ performance gap into an apparent statistical tie until it was found and corrected, and a fixed token budget that materially depressed the observed score of a verbose reasoning model that is competitive on the problems it answered with code. At local-evaluation scale, serving feasibility and harness validity dominate model effects; we argue these deserve first-class attention alongside leaderboard numbers.

1 Motivation

Running capable coding agents locally and privately, on hardware that sits under a desk rather than in a datacenter, is an increasingly common goal. The DGX Spark—GB10 silicon, 128 GB of unified LPDDR5x memory—is a plausible box for it, and 120B-class open-weight coding models are now available in quantized forms that nominally fit. What is under-documented is the gap between “nominally fits” and “actually serves and can be measured honestly.” This study answers a narrow, defensible question rather than a broad comparative one:

What serving, measurement, and evaluation constraints arise when three pinned open-weight model configurations are run locally on one DGX Spark, and what does that imply for reproducible coding-agent benchmark design?

We make three contributions. First, a **serving-feasibility matrix** for GB10/`sm_121a`: which model/quantization/runtime combinations start successfully, and why the others do not. Second, a **three-layer local evaluation harness** with complete raw-artifact retention and one-command deterministic rescoring, so every number in this paper regenerates from committed outputs without re-serving a model. Third, a set of **measurement lessons**—most sharply, a documented case where one missing specification file changed a headline conclusion, and a fixed token budget that depressed a reasoning model’s score in a way that looked like a capability gap but was an output-length artifact.

We state the framing plainly because it governs how every result should be read. Model identity on this box is fully confounded with serving runtime, quantization, and agent scaffold, and that confound is a consequence of hardware, not a design choice

we could relax. The closest available model-versus-model comparison is the quality axis between the two models that share a runtime (gpt-oss and Qwen, both on vLLM); even this is not a clean contrast, since the two still differ in size, quantization (MXFP4 vs. bf16), architecture, and tool/reasoning handling, so we frame it as the closest available comparison rather than a controlled one. Everything that ties a model to a distinct runtime is a configuration comparison and is labeled as such at the point of every affected number.

2 Related Work

Coding-agent benchmarks. SWE-bench [1] and its human-filtered SWE-bench Verified subset established repository-level, execution-graded bug repair as the dominant agentic-coding measure; a patch is scored by whether the repository’s own `FAIL_TO_PASS` and `PASS_TO_PASS` test suites pass after it is applied. We use SWE-bench Verified as our Layer 1 anchor. LiveCodeBench [2] provides time-windowed, single-shot competitive-programming problems with hidden tests, designed to control contamination by selecting problems released after a model’s training cutoff; we use it as our Layer 3 single-shot code-generation axis. Because our fixed pre-2024-06 window predates the training cutoffs of all three models, we describe it as a contamination-possible historical window rather than a contamination-free one. Contamination-hardened and capability-broadened successors such as SWE-bench Pro [5] and cloud-sandboxed agentic suites exist, but they are x86- or cloud-bound and thus off-threshold for a local, on-box study; we treat them as related and future work rather than running them. Our Layer 2—building a full-stack application from a frozen specification and scoring it against an automated HTTP contract—occupies a long-horizon “construct from nothing” axis that few coding-agent benchmarks test.

Local serving of MoE models. vLLM [3] and TensorRT-LLM [4] are the two OpenAI-compatible serving stacks we use. Sparse MoE architectures [6] decouple total parameters from per-token active parameters, which is what makes 120B-class models tractable on a 128 GB box; the three models here span an active-parameter fraction (active divided by total parameters) of 4.4% to 10.8%. Low-bit weight formats—OCP microscaling MXFP4 [13] and NVIDIA’s distinct NVFP4 [14]—shrink the memory footprint further. A recurring and under-reported subtlety, which our results make concrete, is that in the pinned runtime builds we tested on con-

sumer Blackwell (`sm_121a`), FP4 weights were served through weight-only kernels rather than an FP4 *compute* path, so the realized FP4 benefit was storage and bandwidth, not tensor-core arithmetic. This is a property of the software path we observed on these runtime builds, not a statement about the GB10 tensor cores, which do support FP4 arithmetic.

Evaluation methodology. Our statistical treatment uses Wilson score intervals [7] for binomial rates and McNemar’s test [8] for paired comparisons on shared task sets. The broader point we contribute to—that measured scores are highly sensitive to harness details, sample size, and token budgets—is consistent with the wider reproducibility literature, but we ground it in specific, reproducible failures on this hardware rather than in the abstract.

3 Experimental Setup

3.1 Hardware

The box is a single DGX Spark [15]: GB10, aarch64, consumer Grace Blackwell, compute capability `sm_121a`, with 128 GB of unified LPDDR5x shared between CPU and GPU at roughly 273 GB/s—about 6× lower than discrete Blackwell’s ~1.7 TB/s. Decode is therefore memory-bandwidth-bound, and this shapes several results. Software: driver 580.159.03, CUDA 13.0, device reported as `NVIDIA GB10`. One large model is served at a time.

A hardware detail that matters for measurement: because memory is unified, `nvidia-smi` runs correctly but its `memory.*` fields return `Not Supported/[N/A]` by design—there is no separate VRAM pool to count. The model’s memory footprint is therefore read from `/proc/meminfo` as host RAM and must not be interpreted as dedicated GPU memory. This is a property of the platform, not a broken tool, and we disclose it wherever a memory figure appears.

3.2 Serving and agent

Models are served behind OpenAI-compatible endpoints: vLLM on port 8000 and TensorRT-LLM on port 8355. Each serve passes a correctness gate (a fixed prompt must produce coherent, non-repeating output; a failed gate aborts the run rather than recording garbage metrics) and writes a per-serve reproducibility manifest. Container image digests are pinned: vLLM `nvcr.io/nvidia/vllm:26.02-py3` (vLLM 0.15.1) and TensorRT-LLM `nvcr.io/nvidia/tensorrt-llm:release:1.3.0rc9`. The agentic layers are driven by OpenCode [9] 1.17.9 at an identical version and configuration for every

model; Layer 3 uses no agent and calls the endpoint directly.

3.3 The three configurations

Table 1 lists the three arms. All are sparse MoE models with an active-parameter fraction of 4.4% (gpt-oss), 10.0% (nemotron), and 10.8% (qwen). Serving settings cannot be equalized across all three arms, and we do not claim they are. Random seed (0) and sampling temperature are held identical across all three arms at the serving layer, but only the two vLLM arms (gpt-oss, qwen) additionally share KV-cache dtype policy, `max_num_seqs`, and the OpenCode and vLLM versions. The nemotron arm runs on TensorRT-LLM and necessarily differs: FP8 KV cache, a serving batch of 8 (against `max_num_seqs` 1 for the vLLM arms), distinct reasoning/tool parsers, and—at Layer 1/2—MTP speculative decoding; its Layer 3 generation additionally ran 8-way concurrently rather than single-stream. We therefore report per-configuration serving settings (Appendix A) rather than a single equalized protocol, and label every cross-runtime comparison as a configuration comparison. What no arm equalizes is disclosed throughout: quantization, serving runtime, and `max_model_len`/memory-utilization under the 128 GB ceiling.

A candidate fourth model, DeepSeek-Coder-V2-Lite, was evaluated and retired before any timed run. It generates correct code but, in the tested configuration, did not *autonomously* call tools—in an OpenCode smoke test it claimed success while writing nothing to disk—so it could not drive the agent loop. We introduced an explicit autonomous-tool-use gate as a precondition for the agentic layers: a model must return a well-formed `tool_calls` response under `tool_choice:"auto"` (not prose, not hallucinated tool-output tokens) *and* actually create a file on disk. gpt-oss-120b passed this gate and replaced DeepSeek. The case is itself a finding: agentic competence is not the same property as code-generation competence, and a benchmark that conflates them will mis-rank models.

3.4 The headline systems result: serving feasibility

The single most practitioner-useful finding is negative and structural. **No model in our set serves on both runtimes on GB10/sm_121a.** Nemotron is TensorRT-LLM-only; vLLM 0.15.1 as shipped in NGC 26.02-py3 rejects its checkpoint at configuration validation, before any weights load.

The checkpoint declares `quant_algo=MIXED_PRECISION` (FP8 on the Mamba mixer projections, NVFP4 on the MoE experts, per layer, produced by ModelOpt 0.43.0.dev63), and this vLLM build’s ModelOpt loader accepts only a single uniform algorithm from a fixed whitelist. The failure is a `pydantic` validation error raised in ~ 6 s during engine-config construction, not an out-of-memory condition or a runtime crash. We scope this claim precisely: it is a property of *vLLM 0.15.1 as shipped in NGC 26.02-py3*, the stack we pinned for driver compatibility; we did not test a newer vLLM build or the current NVIDIA-validated DGX Spark recipe, so the claim is scoped to the pinned stack, not to vLLM in general. A circulating workaround forces the top-level algorithm to NVFP4; we deliberately did not use it, because the Mamba mixer projections are genuinely FP8 in the weights and forcing NVFP4 risks silently miscasting them—a model-under-test that might load but generate subtly wrong output is worse than one that cleanly refuses.

Conversely, the two vLLM models do not serve on TensorRT-LLM 1.3.0rc9 on this silicon: bf16 MoE has no working `sm_121a` kernel, and MXFP4 MoE autotunes but then deadlocks at the server-to-executor handoff. The consequence is that a same-model, cross-runtime “bridge” run—the clean way to isolate a runtime effect from a model effect—is impossible on this box with the pinned runtimes. This is exactly why we treat the arms as configurations and make no controlled cross-runtime speed or energy claim.

4 Evaluation Design

The three layers (Table 2) probe orthogonal axes of “coding ability.” L1 and L2 are agentic multi-turn tool use driven through OpenCode; L3 is single-shot generation against the endpoint. Keeping them separate is deliberate: it lets “code generation is not the same as agentic tool use” be an empirical observation rather than an assumption.

Layer 1 runs on a 29-task ARM64-buildable subset of SWE-bench Verified, and we are explicit that this is *not* SWE-bench Verified performance and *not* a random sample of the full 500-instance set. The official task images are x86-only. Rather than probe all 500 instances, we drew a stratified sample of up to four instances from each of the 12 SWE-bench Verified repositories (43 instances in total), built each one’s environment natively for aarch64, and kept the 29 whose environment builds and whose gold patch resolves on ARM64. This is an empirical buildability probe over a stratified sample—not the full buildable

Table 1: The three deployment configurations. Model identity is coupled to quantization, runtime, and scaffold by hardware necessity, and each coupling is disclosed. No model in the set serves on both runtimes on this box.

Model	Architecture	Total / Active	Quant	Serves on	Why locked to one runtime
gpt-oss-120b	MoE (128 exp, top-4)	116.8B / 5.1B	MXFP4	vLLM only	MXFP4 MoE autotunes on TRT-LLM but deadlocks at the executor handoff
nemotron-3-super	MoE-hybrid Mamba/attn	~120B / ~12B	NVFP4 (mixed)	TRT-LLM only	vLLM 0.15.1 (26.02-py3) rejects its <code>MIXED_PRECISION</code> checkpoint
qwen3-coder-30b	MoE (128 exp, top-8)	30.5B / 3.3B	bf16	vLLM only	bf16 MoE has no working <code>sm_121a</code> TRT-LLM kernel

Table 2: The three measurement layers. L1/L2 are agentic (multi-turn tool use); L3 is single-shot. Scoring is fully automated and mechanical at every layer—no human scorers and no LLM judge.

Layer	What	Metric
L1	SWE-bench Verified repo bug-fix, agentic	resolved pass@1 on a disclosed 29-task ARM64-buildable subset
L2	Full-stack app build from a frozen spec, agentic	automated API acceptance-check fraction ($k/29$)
L3	LiveCodeBench single-shot generation	pass@1 with Wilson 95% CI ($n=512$)

set, and not a curated difficulty sample; two of the 12 probed repositories (xarray, scikit-learn) failed to build on aarch64 and contribute no tasks. The surviving subset is biased toward pure-Python repositories (flask, requests, pylint, pytest, sympy, sphinx, and similar) and excludes anything needing x86-only wheels or heavy C/Fortran toolchains—a coverage limitation we disclose. Because all three configurations run the identical arm64-built tasks, any ARM-versus-x86 environment difference cancels in the relative comparison. The subset, the 43-instance probe list, and their checksums are pinned in the provenance record.

Layer 2 has each model build the same realistic full-stack application (“TaskFlow Local”) from an identical frozen specification, across a Node.js and a Python/FastAPI backend track. Scoring fires 29 equally weighted, fully automated HTTP assertions against a frozen API contract after booting the built app, reported as an acceptance-check fraction $k/29$. This is an API-acceptance signal, not a holistic app-quality score: it measures whether the app boots and satisfies the communicated contract, not user experience, design, or maintainability. We report a

Bernoulli working-app rate—a run counts as a working app if the built app boots and passes at least half of the 29 checks ($k/29 \geq 0.5$)—with a Wilson interval as the headline, never a single- N mean rubric score. We label the 0.5 cutoff exploratory, and note that the 29 checks are stateful and mutually dependent (later checks reuse state created by earlier ones) rather than independent Bernoulli trials.

Layer 3 runs LiveCodeBench single-shot on a fixed pre-2024-06 window of 512 problems, scored as pass@1 with a Wilson 95% interval. A shared maximum generation budget of 8192 tokens is held identical across all three models as a fairness control—a decision that turns out to be consequential (Section 5.3).

Scoring and fairness. Scoring is deterministic given the retained artifacts, so runs can be rescored from committed outputs without re-serving. Same spec, prompt, hardware, time box, and harness for every run; all serving-setting differences that could not be equalized are disclosed; the maintainer’s tooling never edits model-generated application code, which lives in a separate repository.

5 Results

Figure 2 shows success by layer with Wilson intervals. On raw scores the ordering is $\text{gpt-oss} \geq \text{qwen} \geq \text{nemotron}$, but the L1 and L2 gaps are statistically weak and the L3 nemotron gap is budget-confounded. We report the numbers and then explain, layer by layer, why the durable finding is about how one measures rather than which model wins.

5.1 Layer 1: SWE-bench Verified ARM64 subset

Table 3 gives the Layer 1 rates. Two points about denominators. First, the primary rate we report is the raw, intention-to-treat $6/29 = 20.7\%$: it

Table 3: Layer 1, 29-task ARM64 subset, single-attempt pass@1. The raw $k/29$ rate is primary (intention-to-treat); the “model-valid” rate, which excludes operational failures (a watchdog kill and eight clone/DNS/guard-drops that were never a measured attempt) from the denominator, is reported only as a sensitivity analysis.

Model	Resolved (raw)	Model-valid rate	Notes
gpt-oss-120b	37.9% (11/29)	37.9% (11/29)	no infra failures
qwen3-coder-30b	24.1% (7/29)	24.1% (7/29)	no infra failures
nemotron-super	20.7% (6/29)	30.0% (6/20)	1 watchdog kill; 8 excluded

charges every operational failure (a watchdog kill and eight environment failures—clone/DNS/guard drops—that were never a measured model attempt) against the configuration. As a sensitivity analysis we also report a model-valid rate that excludes those nine non-attempts, $6/20 = 30.0\%$; we flag it as a sensitivity rather than the headline because excluding failures from only one arm can introduce selection bias, and watchdog/runtime failures may themselves be configuration-dependent. Second, pairwise McNemar tests on the shared 29 tasks find no significant separation for any pair, before or after a Holm correction across the three comparisons: gpt-oss vs. nemotron $p = 0.125$ (Holm-adjusted 0.375), gpt-oss vs. qwen $p = 0.289$ (adjusted 0.578), and nemotron vs. qwen $p = 1.0$ (adjusted 1.0). At this sample size the layer cannot rank these models; we read this as a failure to reject, not as evidence of equivalence.

A second axis matters beyond “did it fix the bug.” SWE-bench also records whether a patch broke something that previously passed (PASS_TO_PASS). Among patches that applied and were evaluated, gpt-oss regresses at least one previously-passing test 33.3% of the time (8/24), qwen 31.8% (7/22), and nemotron 21.4% (3/14). The severity diverges far more than the rate: gpt-oss broke 1078 PASS_TO_PASS tests to nemotron’s 3, because a few catastrophic patches (for instance, breaking an import so a whole module’s suite fails) dominate the count. Figure 3 shows the trade-off: gpt-oss’s aggressive edits resolve the most and regress the most, while nemotron’s fewer, smaller edits rarely regress but rarely resolve. This characterizes failure quality among the non-resolving runs, not the winners, and the small denominators keep it directional.

5.2 Layer 2: app-build acceptance, and a harness bug that changed the conclusion

The first Layer 2 sweep was produced under a harness-validity error we label C1. The frozen API contract that the prompt requires the model to satisfy had been dropped from the workspace by

an app-repository tag-lineage break: four of the 29 checks were structurally unreachable, and every model was graded against a contract it could not see. After restoring the contract (app-repo tag **baseline-v7**) and re-running the vLLM models, the picture changed decisively.

Table 4 reports the corrected sweep. Under the bug, gpt-oss (node acceptance 0.252) and qwen (0.155) looked indistinguishable—the working-app Wilson intervals overlapped, and the honest read was “too close to call.” Corrected, gpt-oss (node acceptance 0.724) substantially outperforms qwen (0.178) on the node track: 80% working apps (16/20, Wilson 95% [58.4, 91.9]%) against 15% (3/20, [5.2, 36.0]%), and the working-app intervals no longer overlap. We report this descriptively: the 20 runs per cell are repeated trials of a single application specification, not 20 independent tasks, and the harness applies no inferential test at Layer 2, so we make no population-level significance claim. (Were the runs treated as independent, a Fisher exact test on the node working-app counts would give $p \approx 9 \times 10^{-5}$; we report it only to bound the effect, not as a validated result.) The before/after are also two separate stochastic sweeps rather than a contemporaneous controlled A/B, so we describe the change as observed, not as a clean causal isolation. Crucially, the contract helped *only where a model could already build a booting app*: gpt-oss’s node score rose $2.9\times$ (the four contract-only checks went from 0 to 48 passes), while qwen and both python tracks, which are floored for other reasons, barely moved. A single missing specification file had compressed a roughly $4\times$ performance gap on this task into an apparent statistical tie. Figure 4 shows the before/after.

The asterisk on nemotron marks an asymmetry we disclose explicitly. Nemotron was *not* rerun with the visible contract, so its row is a pre-C1 measurement—retained and marked, but not comparable to the corrected vLLM rows. We did not rerun it because at Layer 2 it almost never produced a working backend (node acceptance mean 0.009, zero working apps) and the four contract-only checks require a booted backend to be reachable; we judged a rerun a poor use of

Table 4: Layer 2 after the C1 contract fix (contract-visible **baseline-v7**) for the two vLLM arms. Working-app rate counts a run that boots and passes at least half of the 29 checks. The nemotron row (starred) is a pre-C1 measurement that was *not* rerun; it is shown for reference only and is not comparable to the corrected rows.

Model	node $k/29$	node working	python $k/29$
gpt-oss-120b	0.724	16/20 = 80%	0.069 (0/8)
qwen3-coder-30b	0.178	3/20 = 15%	0.004 (0/8)
nemotron-super*	0.009 (N=4)	0/4	0.000 (N=4)

fragile TensorRT-LLM compute, but we do not claim to have measured that the contract would have no effect on it. For the vLLM models, results are dual-reported as $k/29$ and a C1-rescored $k/25$ over the 25 reachable checks (gpt-oss node 0.744, qwen 0.204). Two caveats survive the correction: the python track floors for both vLLM models, a genuine node-versus-python asymmetry rather than a contract artifact; and the layer remains an API-acceptance signal, not a full app-quality score.

5.3 Layer 3: LiveCodeBench, and a token budget that manufactured a last place

Table 5 gives raw Layer 3 pass@1, where nemotron finishes last. That last place is substantially driven by output truncation under the shared 8192-token budget: nemotron’s verbose reasoning left 185 of 512 problems with no extractable code (143 empty outputs plus 42 with output but no code block), all scored as fails, against 7/512 for gpt-oss and 10/512 for qwen. Whether a larger budget would recover correct answers on those 185 problems was not measured, so we report the truncation itself rather than an imputed higher score. The truncation is difficulty-correlated—no-code rate 14.3% on easy, 34.3% on medium, 71.5% on hard problems—so a rate conditional on only the answered subset must not be compared against the other models’ full-512 rates, a selection-bias trap we flag (finding C2). As a conditional analysis, on the 327 problems nemotron answered with code—an easier subset selected by nemotron’s own behavior, so this is accuracy *conditional on emitting extractable code* and is not comparable to the full-set rates—nemotron scores 96.0% (314/327), gpt-oss 95.4% (312/327), and qwen 82.3% (269/327). The primary Layer 3 score for every model remains the full-set rate in Table 5. Figure 5 shows both halves.

The methodology finding is general: a single-shot, fixed-budget benchmark systematically penalizes verbose reasoning models, and the penalty concentrates on hard problems. We keep 8192 tokens as the pri-

mary, controlled number—holding the budget constant is the fairness control—and treat the truncation itself as a reportable result rather than tuning it away.

6 Measurement Pitfalls

Several pitfalls surfaced only because we retained raw artifacts and re-derived every table. They generalize beyond this box.

Unified memory is not VRAM. As noted, `nvidia-smi` reports [N/A] for device memory on GB10 by design; the footprint comes from `/proc/meminfo` and is whole-system, not dedicated GPU memory.

Tokens-per-joule is prompt-token-dominated. The agent loop re-sends context every turn, so prompt tokens are a median 98.7% of the total. A naive `total_tokens / energy` therefore measures re-submitted context, not decode work; we report energy per *generated* token alongside it.

Restart-truncated energy windows. Watchdog restarts reuse a run identifier and overwrite the metric window, so a crash-prone run’s energy can cover only its final segment. The gpt-oss Layer 3 energy cell is a known-broken capture (roughly three orders of magnitude low) and is marked invalid in the ledger rather than published—the pass@1 result is unaffected because scoring is independent of the metrics window.

Robust energy statistics. Arithmetic-mean Layer 1 energy is dominated by a single ~ 13 h nemotron operational runaway (~ 655 kJ). We report median and IQR with an operational-failure sensitivity, never a bare mean.

Telemetry asymmetry. vLLM exposes Prometheus histograms (gpt-oss ~ 26.9 tok/s, qwen ~ 19.7 tok/s at Layer 1); TensorRT-LLM exposes none, so nemotron has hardware energy but no decode-throughput or TTFT figures. Those cells are blank by design, not missing data, and we do not impute them.

Table 5: Layer 3 LiveCodeBench pass@1 (pre-2024-06 window, $n=512$). Nemotron’s raw rate is depressed by output truncation under the shared 8192-token budget (see text). The no-code column counts problems with empty output or no extractable code block, all scored as fails.

Model	pass@1	Wilson 95% CI	no-code (of 512)
gpt-oss-120b	89.3%	[86.3, 91.7]	7 (1.4%)
qwen3-coder-30b	68.2%	[64.0, 72.1]	10 (2.0%)
nemotron-super	61.3%	[57.0, 65.4]	185 (36.1%)

7 Operational Findings

A few operational notes for anyone reproducing this. TensorRT-LLM 1.3.0rc9 on GB10 is fragile for non-NVIDIA MoE models—the FP4-native NVIDIA reasoning model serves, the others deadlock. A degraded-box state is real: after a 26-hour vLLM run the GPU was left in a state where TensorRT-LLM nemotron wedged after one generation (warm-up 656 s); a reboot restored normal operation (warm-up back to 8 s), so the box should be reset between heavy long-lived serves. A non-streaming request path wedged nemotron mid-generation, and switching the LiveCodeBench runner to stream (the path OpenCode already used) fixed it. Throughput is bandwidth-bound: nemotron’s single-stream Layer 3 would take ~ 130 h, so we ran it 8-way concurrently (~ 9.5 h). We did not run a 1-way-versus-8-way sensitivity check, so we disclose this concurrency difference as a further configuration confound rather than assert it is quality-neutral; pass@1 is scored per problem and independent of the metrics window, and TensorRT-LLM exposes no throughput/energy histograms to lose. Finally, several silent-failure traps were guarded: reasoning models return `content=None` on truncation; a resume filter re-generated empty outputs indefinitely; and a watchdog nearly false-stalled on a frozen log mtime and now uses the TensorRT-LLM iteration counter as ground truth. Figures 6 and 7 summarize the efficiency reality, with the caveats above embedded.

8 Reproducibility

Every run keeps its summaries, manifests, and clock records under `results/raw/<run-id>/`; bulky per-run time series are regenerable and excluded from version control, while summaries and manifests are committed. Scoring is fully automated and mechanical—no human scorers, no LLM judge, at any layer—so the aggregation and robust-summary scripts regenerate their numeric tables deterministically from retained outputs (rendered figures may differ byte-wise across matplotlib/toolchain versions, but the un-

derlying values do not). A single command (`make rebuild`) re-derives every summary table, per-run ledger, figure, and statistic from the raw runs without re-serving any model, and the statistics and robust-summary self-tests pass. Container image digests, upstream dataset snapshot revisions (SWE-bench Verified, LiveCodeBench), the exact HuggingFace revisions of all three model checkpoints, the Layer 1 subset checksums, and the frozen Layer 3 window are pinned in an immutable provenance record. Model-generated application code lives only in the companion app repository and is never edited by the maintainer’s tooling; the coding assistant used to build the harness and write this report is not part of what was measured.

9 Threats to Validity and Limitations

Single box, three configurations. Model identity is fully coupled to serving runtime, precision, and scaffold. Conclusions are restricted to these configurations on this machine, and serving incompatibility is a systems finding, not proof of intrinsic quality.

Small and unequal samples. L1 and L3 are single-attempt per cell; L2 is $N = 20/8$ (node/python) for the vLLM models against $N = 4$ for nemotron. All quality claims are descriptive, and the McNemar results show L1 cannot rank these models at this size.

Protocol changed mid-collection. The C1 contract fix and rubric-denominator changes mean some early runs are exploratory; confirmatory claims draw on the post-fix `baseline-v7` sweep.

L1 is a disclosed subset, 29 tasks kept from a 43-instance stratified (up-to-four-per-repo) buildability probe, not SWE-bench Verified, not the full ARM64-buildable set, and not a random sample of the 500 instances; absolute SWE-bench numbers should not be extrapolated from it.

L3 is a fixed historical window where contamination is possible: equal chronological eligibility is not equal contamination, the set is not exposure-

balanced, and it is not leaderboard-comparable.

Untested newer stacks. The nemotron/vLLM incompatibility is scoped to the pinned vLLM 0.15.1 (NGC 26.02-py3); a newer vLLM or the current NVIDIA-validated DGX Spark recipe was not tested and could change that specific result.

Isolation. Generated code is executed without OS-level sandboxing on this kernel (bubblewrap unavailable); publication-grade runs should use a no-network container or VM.

10 Conclusion

On a single DGX Spark the story that survives scrutiny is not “model X wins.” It is that at local-evaluation scale, serving feasibility and harness validity dominate model effects. No model in our set serves on both runtimes on this box; a single missing specification file changed the Layer 2 conclusion until it was found; and a fixed token budget substantially depressed the Layer 3 score of a verbose reasoning model that, on the problems it answered with code, scored comparably to the field (whether a larger budget would recover its full-set score is untested). The practitioner takeaways are concrete. Of the configurations measured here, gpt-oss-120b on vLLM is the strongest local agentic setup (L1 37.9%, L2 node 80% working apps, L3 89.3%); nemotron serves only on TensorRT-LLM and its Layer 3 score is heavily affected by the shared token budget; and honest local benchmarking depends far more on a failure taxonomy, robust energy statistics, and deterministic rescoring than on producing one more leaderboard number.

Data and Code Availability

The harness, frozen specification, raw run summaries, aggregation and statistics code, figures, and provenance pins are in the public repository at <https://github.com/mani-mal/spark-coder-bench>. Every table and figure in this paper maps to a committed source CSV and regenerates via `make rebuild`. The model-generated application code lives in the separate public companion repository <https://github.com/mani-mal/taskflow-local-app-benchmark>; all reported Layer 2 runs start from its frozen tag `baseline-v7` (commit 5b0b85b).

References

- [1] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? *ICLR*, 2024. arXiv:2310.06770.
- [2] N. Jain, K. Han, A. Gu, W.-D. Li, F. Yan, T. Zhang, S. Wang, A. Solar-Lezama, K. Sen, and I. Stoica. LiveCodeBench: Holistic and Contamination-Free Evaluation of Large Language Models for Code. arXiv:2403.07974, 2024.
- [3] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica. Efficient Memory Management for Large Language Model Serving with PagedAttention. *SOSP*, 2023. arXiv:2309.06180.
- [4] NVIDIA. TensorRT-LLM: A TensorRT toolbox for optimized large language model inference. <https://github.com/NVIDIA/TensorRT-LLM>. Container image `release:1.3.0rc9`.
- [5] Scale AI. SWE-bench Pro: A contamination-resistant, enterprise-grade software engineering benchmark. 2025. https://github.com/scale-ai/SWE-bench_Pro-os.
- [6] W. Fedus, B. Zoph, and N. Shazeer. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. *JMLR*, 2022. arXiv:2101.03961.
- [7] E. B. Wilson. Probable Inference, the Law of Succession, and Statistical Inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927.
- [8] Q. McNemar. Note on the Sampling Error of the Difference Between Correlated Proportions or Percentages. *Psychometrika*, 12(2):153–157, 1947.
- [9] OpenCode. OpenCode: an open-source terminal-based coding agent. Version 1.17.9. <https://opencode.ai>.
- [10] OpenAI. gpt-oss-120b & gpt-oss-20b Model Card. arXiv:2508.10925, 2025. Card: <https://huggingface.co/openai/gpt-oss-120b>.
- [11] Qwen Team. Qwen3-Coder-30B-A3B-Instruct model card. <https://huggingface.co/Qwen/Qwen3-Coder-30B-A3B-Instruct>, 2025.
- [12] NVIDIA. NVIDIA-Nemotron-3-Super-120B-A12B model card. <https://huggingface.co>

/nvidia/NVIDIA-Nemotron-3-Super-120B-A
12B-NVFP4, 2026.

- [13] Open Compute Project. OCP Microscaling Formats (MX) Specification, Version 1.0. <https://www.opencompute.org/documents/ocp-microscaling-formats-mx-v1-0-spec-final-pdf>, 2023.
- [14] NVIDIA. Introducing NVFP4 for Efficient and Accurate Low-Precision Inference. NVIDIA Technical Blog, 2025. <https://developer.nvidia.com/blog/introducing-nvfp4-for-efficient-and-accurate-low-precision-inference/>.
- [15] NVIDIA. NVIDIA DGX Spark (GB10 Grace Blackwell): product specifications. <https://www.nvidia.com/en-us/products/workstations/dgx-spark/>, 2025.

DGX Spark Local Coding-Model Benchmark — Architecture (3 models × 3 layers, one model served at a time on GB10 / 128 GB unified memory)




Figure 1: End-to-end architecture of one benchmark run on a single DGX Spark. Sweep drivers read a single model registry and hand each arm its model plus serving profile; the selected model comes up behind an OpenAI-compatible endpoint (vLLM on :8000 or TensorRT-LLM on :8355) after a sanity gate and a per-serve reproducibility manifest; OpenCode drives the two agentic layers (L1, L2) while L3 calls the endpoint directly; synchronized collectors capture hardware energy/memory, vLLM Prometheus throughput, and agent accounting into per-run summaries that aggregate into the committed CSVs behind every table here.




Figure 2: Success by layer for the three configurations, each with a Wilson 95% interval: L1 resolved pass@1, L2 node working-app rate, L3 pass@1. The raw ordering is $\text{gpt-oss} \geq \text{qwen} \geq \text{nemotron}$, but the L1/L2 separations are weak and the L3 nemotron result is a token-budget artifact (Figure 5).




Figure 3: Layer 1 collateral damage, as a fraction of applied patches. gpt-oss both resolves the most and regresses the most; nemotron regresses least but resolves little. Aggressive editing buys resolutions at the cost of regressions.




Figure 4: The C1 ablation. Under the harness bug (contract invisible), gpt-oss (6/20) and qwen (4/20) working-app rates overlap. With the contract restored, gpt-oss jumps to 16/20 while qwen stays flat and the intervals separate. One missing spec file changed the observed conclusion (before/after are separate stochastic sweeps, not a contemporaneous A/B).




Figure 5: Left: nemotron's no-code (truncation) rate rises with difficulty (easy 14%, medium 34%, hard 72%). Right: on the 327 problems it answers with code—a conditional, easier subset, not comparable to full-set rates—nemotron scores 96.0%, gpt-oss 95.4%, and qwen 82.3%. Much of nemotron's last place in Table 5 reflects truncation under the fixed budget rather than incorrect answers.





Figure 6: Median wall-clock per task, agentic layers. Nemotron is roughly 2× slower than the vLLM models at L1 (3.9 vs 1.4–1.9 min) and slowest at L2 (16.1 min); the median is shown because the L1 mean is skewed by one ~13 h runaway.



Lower is better. Blank = 0 successes (nemotron L2). GPU-reported energy (nvidia-smi power.draw, GPU package only). L1 excludes operational-failure energy; L2 is Node-track only. Runtime-confounded (nemotron on TRLLM).

Figure 7: Energy per *successful* task (GPU-reported `nvidia-smi power.draw`, GPU package only, no idle subtraction; runtime-confounded, nemotron on TensorRT-LLM). At L1, nemotron costs ~ 33.8 kJ per solve—over its 20 model-valid energy windows, excluding one ~ 13 h watchdog runaway (~ 655 kJ, 76% of its L1 energy)—to gpt-oss’s 12.8 kJ ($\sim 2.6\times$); including the runaway inflates it to 143 kJ. At L2, energy per working app is Node-track only (all working apps are Node): qwen 155.3 kJ to gpt-oss’s 22.1 kJ; pooling in the Python track, which produced no working apps, inflates these to 211 and 32.8 kJ.

A Serving Commands and Configurations

Both runtimes are launched from a single sequential driver that writes a per-serve reproducibility manifest and passes a coherence sanity gate before any timed run. The vLLM invocation (Qwen shown; per-model values come from the `infra/models.json` registry, and `-max-num-seqs 1` is pinned across all models for clean per-task latency):

```
docker run -d --name vllm-server --gpus all --ipc=host \
  --network=host --ulimit memlock=-1 --ulimit stack=67108864 \
  ↪ \
  -v $HF_HOME:/root/.cache/huggingface \
  nvcr.io/nvidia/vllm:26.02-py3 \
  vllm serve Qwen/Qwen3-Coder-30B-A3B-Instruct \
  --host 0.0.0.0 --port 8000 --served-model-name qwen3- \
  ↪ coder-30b \
  --gpu-memory-utilization 0.80 --max-model-len 65536 \
  --max-num-seqs 1 --max-num-batched-tokens 16384 \
  --dtype auto --trust-remote-code --seed 0 \
  --enable-auto-tool-choice # + per-profile tool-call \
  ↪ parser
```

TensorRT-LLM serves Nemotron on port 8355 with an `extra_llm_api_options` YAML. The 1.3.0rc9 image is required: 1.2.1’s `nemotron_h` loader fell back to a regular init and double-allocated to ~115 GB, OOM-ing the host.

```
trtllm-serve nvidia/NVIDIA-Nemotron-3-Super-120B-A12B-NVFP4 \
  --host 0.0.0.0 --port 8355 \
  --max_batch_size 8 --tp_size 1 --ep_size 1 \
  --max_num_tokens 8192 --max_seq_len 1048576 \
  --extra_llm_api_options /config/nemotron-super.yaml \
  --reasoning_parser nano-v3 --tool_parser qwen3_coder
```

The run-time serving values above come from the committed per-serve manifests (`max_batch_size 8`, `max_seq_len 1048576`—the 1M sequence length is feasible because the KV cache is FP8). The Layer 3 profile is identical except that MTP speculative decoding is disabled (a KV-manager wedge under MTP, documented in the findings), so the `speculative_config` block below applies to Layers 1–2 only.

The YAML encodes the 128 GB unified-memory strategy for this hybrid Mamba/attention MoE (FP8 KV cache, float16 SSM cache with stochastic rounding, chunked prefill):

```
kv_cache_config:
  dtype: fp8
  enable_block_reuse: false
  free_gpu_memory_fraction: 0.9
  mamba_ssm_cache_dtype: float16
  mamba_ssm_stochastic_rounding: true
moe_config:
  backend: CUTLASS
cuda_graph_config:
  enable_padding: true
  max_batch_size: 8
enable_chunked_prefill: true
speculative_config:
  decoding_type: MTP
  num_nextn_predict_layers: 3
```

B The Nemotron vLLM Rejection

This is the headline serving-feasibility result of Section 3.4. The checkpoint’s `hf_quant_config.json` (produced by `modelopt 0.43.0.dev63`) declares a *mixed-precision* format—FP8 on the Mamba mixer projections, NVFP4 on the MoE experts, specified per layer:

```
{ "quantization": {
  "quant_algo": "MIXED_PRECISION",
  "kv_cache_quant_algo": "FP8",
  "quantized_layers": {
    "backbone.layers.0.mixer.in_proj": {"quant_algo": "FP8"},
    ↪ "backbone.layers.0.mixer.out_proj": {"quant_algo": "FP8"},
    ↪ "backbone.layers.1.mixer.experts.0.up_proj":
      {"quant_algo": "NVFP4", "group_size": 16},
    "...": "~every MoE expert listed, NVFP4 group_size 16"
  } } }
```

vLLM 0.15.1’s ModelOpt loader accepts only a single uniform algorithm and rejects the model during `create_engine_config(...)`—in ~6s, before any weight loads, and before any tool/reasoning-parser flag is reached:

```
pydantic_core._pydantic_core.ValidationError: 1 validation
error for VllmConfig
Value error, ModelOpt currently only supports:
['FP8', 'FP8_PER_CHANNEL_PER_TOKEN', 'FP8_PB_WO', 'NVFP4']
quantizations in vLLM. Please check the
'hf_quant_config.json' file for your model's quant config.
```

No profile or parser change fixes this: the rejected value is the checkpoint’s own format. A circulating workaround forces the top-level field to `NVFP4`; we did not use it, because the Mamba mixer projections are genuinely FP8 in the weights and forcing NVFP4 risks silently miscasting them. The same error against the same model and container is filed upstream at `vllm-project/vllm#37854` and `NVIDIA/dgx-spark-playbooks#77`.

C The Autonomous Tool-Use Gate

Before any timed agentic run, a model must pass a two-part gate; a model that fails is reported as a non-agentic baseline rather than scored as if it tried. (1) A direct request to the served endpoint with `tool_choice:"auto"` and a tool schema must return a non-null, well-formed `tool_calls` object—not prose, and not hallucinated tool-output tokens in the content. (2) In an OpenCode smoke test the model must actually create a file on disk. A model that only acts under `tool_choice:"required"` fails part (1). This

is the gate DeepSeek-Coder-V2-Lite failed—it generated correct code but wrote nothing to disk while reporting success—and that gpt-oss-120b passed before replacing it.

D Layer 2 API Contract

Layer 2 scores 29 equally weighted, fully automated HTTP assertions after booting the built app. The checks and their categories:

```
boot: health_ok
auth: login_admin_ok, login_bad_creds_401,
      me_without_token_401, me_with_token_200,
      login_member_ok
security: user_no_password_field,
          protected_route_no_token_401
rbac: member_create_project_403
projects: admin_create_project, project_list_contains,
          project_get_by_id, project_edit*, project_archive*
tasks: task_create, task_list, task_get,
       task_update_status*, task_filter_status,
       task_search, task_delete
comments: comment_add, comment_list
validation: validation_project_no_name_400,
            validation_task_no_title_400,
            validation_task_bad_priority_400
dashboard: dashboard_summary_keys*
frontend: frontend_build
reproducibility: lockfile_present
```

The four checks marked * are the C1 contract-only assertions of Section 5.2: they pin exact routes and key sets the requirements prose leaves unspecified, so they are reachable only when the frozen `api-contract.md` is in the workspace. Each check is a small deterministic assertion; the RBAC bypass check—which caught apps that let a non-admin member create a project—is representative:

```
# member (non-admin) must be forbidden from creating a
#   ↪ project
r, _ = self.req("POST", C.API + "/projects",
               token=member_token, json={"name": "X"})
self.add("member_create_project_403", "rbac",
        r is not None and r.status_code == 403,
        r.status_code if r else "none")
```

E Layer 1 ARM64-Buildable Subset

The 29 tasks are the SWE-bench Verified instances that built natively on aarch64 and whose gold patch resolved there, drawn from a 43-instance stratified probe (up to four instances from each of the 12 SWE-bench Verified repositories)—an empirical buildability probe, not the full buildable set and not a difficulty sample. Ten of the 12 probed repositories contributed at least one surviving task; xarray and scikit-learn built on neither. The distribution across the surviving repositories:

Repository	N	Repository	N
astropy	4	pylint	3
django	3	pytest	4
matplotlib	2	requests	4
seaborn	2	sphinx	2
flask	1	sympy	4

```
astropy__astropy-12907 astropy__astropy-13033
astropy__astropy-13236 astropy__astropy-13398
django__django-10554 django__django-10880
django__django-10914 matplotlib__matplotlib-13989
matplotlib__matplotlib-14623 mwaskom__seaborn-3069
mwaskom__seaborn-3187 pallets__flask-5014
psf__requests-1142 psf__requests-1724
psf__requests-1766 psf__requests-1921
pylint-dev__pylint-4551 pylint-dev__pylint-4661
pylint-dev__pylint-4970 pytest-dev__pytest-10051
pytest-dev__pytest-10081 pytest-dev__pytest-10356
pytest-dev__pytest-5262 sphinx-doc__sphinx-10449
sphinx-doc__sphinx-10466 sympy__sympy-11618
sympy__sympy-12096 sympy__sympy-12419
sympy__sympy-12481
```

F Per-Model, Per-Layer Resource Metrics

Table 6 gives the consolidated hardware figures behind the efficiency discussion. Nemotron’s decode-throughput and TTFT cells are blank because TensorRT-LLM exposes no Prometheus histograms (Section 6), not because the data is missing. Layer 3 is omitted here: its duration is a whole-run figure rather than per-task, and the gpt-oss L3 metrics window failed to capture the generation (Section 6). Energy and wall-clock are per-task medians; memory is peak unified (system) RAM.

Table 6: Per-model, per-layer resource metrics (medians). Source: `results/summary/perf-resource-summary.csv`.

Layer	Model	N	min/task	kJ/task	GPU util %	GPU W	peak mem (MiB)	decode tok/s	TTFT p50 (s)
L1	gpt-oss-120b	29	1.9	3.6	90.7	30.1	113,059	26.2	0.47
L1	qwen3-coder-30b	29	1.4	2.6	89.3	29.2	110,455	20.1	0.50
L1	nemotron-super	21	3.9	8.4	91.8	38.0	108,517	—	—
L2	gpt-oss-120b	27	11.1	18.8	88.5	30.8	111,132	31.0	0.36
L2	qwen3-coder-30b	28	12.9	21.5	92.7	27.7	110,250	22.8	0.21
L2	nemotron-super	8	16.1	40.6	90.6	45.3	107,594	—	—