

Local Coding-Model Configurations on a Single DGX Spark:

A Systems and Methodology Case Study

Mani Malarvannan

mani.malarvannan@cdw.com

<https://github.com/mani-mal/spark-coder-bench>

Abstract

We report a reproducible case study of serving and evaluating three open-weight Mixture-of-Experts (MoE) coding models locally on a single NVIDIA DGX Spark (GB10, consumer Grace Blackwell, `sm_121a`, 128 GB unified memory): `gpt-oss-120b` [8] (MXFP4, vLLM), `Qwen3-Coder-30B-A3B` [9] (bf16, vLLM), and NVIDIA Nemotron-3-Super-120B-A12B [10] (NVFP4, TensorRT-LLM). This is deliberately *not* a model-ranking paper. On this hardware each model couples to exactly one serving runtime, one quantization, and one agent scaffold, so every score describes a *deployment configuration*, not intrinsic model quality. Our contributions are three: (1) a serving-feasibility matrix for GB10—the headline result is negative and structural, that no model in our set serves on both available runtimes, which makes a matched same-model cross-runtime comparison infeasible on this box; (2) a three-layer local evaluation harness (SWE-bench Verified agentic repair on a disclosed ARM64-buildable subset, a from-scratch full-stack app build scored by an automated API contract, and LiveCodeBench single-shot generation) with full raw-artifact retention and byte-identical deterministic rescore; and (3) a set of measurement lessons drawn from concrete failures. The sharpest lesson is a single harness-validity error that compressed a real $\sim 4\times$ quality gap into a fake statistical tie until it was found and corrected, and a fixed token budget that manufactured a last-place finish for a verbose reasoning model that is competitive on the problems it actually answers. At local-evaluation scale, serving feasibility and harness validity dominate model effects; we argue these deserve first-class attention alongside leaderboard numbers.

1 Motivation

Running capable coding agents locally and privately, on hardware that sits under a desk rather than in a datacenter, is an increasingly common goal. The DGX Spark—

GB10 silicon, 128 GB of unified LPDDR5x memory—is a plausible box for it, and 120B-class open-weight coding models are now available in quantized forms that nominally fit. What is under-documented is the gap between “nominally fits” and “actually serves and can be measured honestly.” This study answers a narrow, defensible question rather than a broad comparative one:

What serving, measurement, and evaluation constraints arise when three pinned open-weight model configurations are run locally on one DGX Spark, and what does that imply for reproducible coding-agent benchmark design?

We make three contributions. First, a **serving-feasibility matrix** for GB10/`sm_121a`: which model/quantization/runtime combinations actually come up, and why the others do not. Second, a **three-layer local evaluation harness** with complete raw-artifact retention and one-command deterministic rescore, so every number in this paper regenerates from committed outputs without re-serving a model. Third, a set of **measurement lessons**—most sharply, a documented case where one missing specification file inverted a headline conclusion, and a fixed token budget that penalized a reasoning model in a way that looked like a capability gap but was not.

We state the framing plainly because it governs how every result should be read. Model identity on this box is fully confounded with serving runtime, quantization, and agent scaffold, and that confound is a consequence of hardware, not a design choice we could relax. The one comparison that is a clean model-versus-model contrast is the quality axis between the two models that share a runtime (`gpt-oss` and `Qwen`, both on vLLM); everything that ties a model to a distinct runtime is a configuration comparison and is labeled as such at the point of every affected number.

2 Related Work

Coding-agent benchmarks. SWE-bench [1] and its human-filtered SWE-bench Verified subset established repository-level, execution-graded bug repair as the dominant agentic-coding measure; a patch is scored by whether the repository’s own `FAIL_TO_PASS` and `PASS_TO_PASS` test suites pass after it is applied. We use SWE-bench Verified as our Layer 1 anchor. LiveCodeBench [2] provides time-windowed, single-shot competitive-programming problems with hidden tests, designed to control contamination by selecting problems released after a model’s training cutoff; we use it as our Layer 3 contamination-aware, code-generation axis. Contamination-hardened and capability-broadened successors such as SWE-bench Pro and cloud-sandboxed agentic suites exist, but they are x86- or cloud-bound and thus off-thesis for a local, on-box study; we treat them as related and future work rather than running them. Our Layer 2—building a full-stack application from a frozen specification and scoring it against an automated HTTP contract—occupies a long-horizon “construct from nothing” axis that few coding-agent benchmarks test.

Local serving of MoE models. vLLM [3] and TensorRT-LLM are the two OpenAI-compatible serving stacks we use. Sparse MoE architectures [4] decouple total parameters from per-token active parameters, which is what makes 120B-class models tractable on a 128 GB box; the three models here span an active/total sparsity range of 4.4% to 10.8%. Low-bit microscaling weight formats (MXFP4, NVFP4) [11] shrink the memory footprint further. A recurring and under-reported subtlety, which our results make concrete, is that on consumer Blackwell (`sm_121a`) there is no native FP4 *compute* path in these runtimes: FP4 weights are served through weight-only kernels, so any FP4 benefit is storage and bandwidth, not tensor-core arithmetic.

Evaluation methodology. Our statistical treatment uses Wilson score intervals [5] for binomial rates and McNemar’s test [6] for paired comparisons on shared task sets. The broader point we contribute to—that measured scores are highly sensitive to harness details, sample size, and token budgets—is consistent with the wider reproducibility literature, but we ground it in specific, reproducible failures on this hardware rather than in the abstract.

3 Experimental Setup

3.1 Hardware

The box is a single DGX Spark: GB10, aarch64, consumer Grace Blackwell, compute capability `sm_121a`, with 128 GB of unified LPDDR5x shared between CPU and GPU at roughly 273 GB/s—about 6× lower than

discrete Blackwell’s ~ 1.7 TB/s. Decode is therefore memory-bandwidth-bound, and this shapes several results. Software: driver 580.159.03, CUDA 13.0, device reported as `NVIDIA GB10`. One large model is served at a time.

A hardware detail that matters for measurement: because memory is unified, `nvidia-smi` runs correctly but its `memory.*` fields return `Not Supported/[N/A]` by design—there is no separate VRAM pool to count. The model’s memory footprint is therefore read from `/proc/meminfo` as host RAM and must not be interpreted as dedicated GPU memory. This is a property of the platform, not a broken tool, and we disclose it wherever a memory figure appears.

3.2 Serving and agent

Models are served behind OpenAI-compatible endpoints: vLLM on port 8000 and TensorRT-LLM on port 8355. Each serve passes a correctness gate (a fixed prompt must produce coherent, non-repeating output; a failed gate aborts the run rather than recording garbage metrics) and writes a per-serve reproducibility manifest. Container image digests are pinned: vLLM `nvcr.io/nvidia/vllm:26.02-py3` (vLLM 0.15.1) and TensorRT-LLM `nvcr.io/nvidia/tensorrt-llm/release:1.3.0rc9`. The agentic layers are driven by OpenCode at an identical version and configuration for every model; Layer 3 uses no agent and calls the endpoint directly.

3.3 The three configurations

Table 1 lists the three arms. All are sparse MoE models spanning an active/total sparsity range of 4.4% (gpt-oss), 10.0% (nemotron), and 10.8% (qwen). We hold identical across arms: random seed, temperature, KV-cache dtype policy, `max_num_seqs`, and the OpenCode and vLLM versions. What we cannot equalize is disclosed: quantization, serving runtime, and `max_model_len`/memory-utilization settings under the 128 GB ceiling.

A candidate fourth model, DeepSeek-Coder-V2-Lite, was evaluated and retired before any timed run. It generates correct code but will not *autonomously* call tools—in an OpenCode smoke test it claimed success while writing nothing to disk—so it cannot drive the agent loop. We introduced an explicit autonomous-tool-use gate as a precondition for the agentic layers: a model must return a well-formed `tool_calls` response under `tool_choice:"auto"` (not prose, not hallucinated tool-output tokens) *and* actually create a file on disk. gpt-oss-120b passed this gate and replaced DeepSeek. The case is itself a finding: agentic competence is not the same property as code-generation competence, and a benchmark that conflates them will mis-rank models.

Table 1: The three deployment configurations. Model identity is coupled to quantization, runtime, and scaffold by hardware necessity, and each coupling is disclosed. No model in the set serves on both runtimes on this box.

Model	Architecture	Total / Active	Quant	Serves on	Why locked to one runtime
gpt-oss-120b	MoE (128 exp, top-4)	116.8B / 5.1B	MXFP4	vLLM only	MXFP4 MoE autotunes on TRT-LLM but deadlocks at the executor handoff
nemotron-3-super	MoE-hybrid Mamba/attn	~120B / ~12B	NVFP4 (mixed)	TRT-LLM only	vLLM 0.15.1 (26.02-py3) rejects its MIXED_PRECISION checkpoint
qwen3-coder-30b	MoE (128 exp, top-8)	30.5B / 3.3B	bf16	vLLM only	bf16 MoE has no working sm_121a TRT-LLM kernel

3.4 The headline systems result: serving feasibility

The single most practitioner-useful finding is negative and structural. **No model in our set serves on both runtimes on GB10/sm_121a.** Nemotron is TensorRT-LLM-only: vLLM 0.15.1 as shipped in NGC 26.02-py3 rejects its checkpoint at configuration validation, before any weights load. The checkpoint declares `quant_algo=MIXED_PRECISION` (FP8 on the Mamba mixer projections, NVFP4 on the MoE experts, per layer, produced by ModelOpt 0.43.0.dev63), and this vLLM build’s ModelOpt loader accepts only a single uniform algorithm from a fixed whitelist. The failure is a `pydantic` validation error raised in ~6s during engine-config construction, not an out-of-memory condition or a runtime crash. We scope this claim precisely: it is a property of *vLLM 0.15.1 as shipped in NGC 26.02-py3*, the stack we pinned for driver compatibility; we did not test a newer vLLM build or the current NVIDIA-validated DGX Spark recipe, so the claim is scoped to the pinned stack, not to vLLM in general. A circulating workaround forces the top-level algorithm to NVFP4; we deliberately did not use it, because the Mamba mixer projections are genuinely FP8 in the weights and forcing NVFP4 risks silently miscasting them—a model-under-test that might load but generate subtly wrong output is worse than one that cleanly refuses.

Conversely, the two vLLM models do not serve on TensorRT-LLM 1.3.0rc9 on this silicon: bf16 MoE has no working sm_121a kernel, and MXFP4 MoE autotunes but then deadlocks at the server-to-executor handoff. The consequence is that a same-model, cross-runtime “bridge” run—the clean way to isolate a runtime effect from a model effect—is impossible on this box with the pinned runtimes. This is exactly why we treat the arms as configurations and make no controlled cross-runtime speed or energy claim.

Table 2: The three measurement layers. L1/L2 are agentic (multi-turn tool use); L3 is single-shot. Scoring is fully automated and mechanical at every layer—no human scorers and no LLM judge.

Layer	What	Metric
L1	SWE-bench Verified repo bug-fix, agentic	resolved pass@1 on a disclosed 29-task ARM64-buildable subset
L2	Full-stack app build from a frozen spec, agentic	automated API acceptance-check fraction ($k/29$)
L3	LiveCodeBench single-shot generation	pass@1 with Wilson 95% CI ($n=512$)

4 Evaluation Design

The three layers (Table 2) probe orthogonal axes of “coding ability.” L1 and L2 are agentic multi-turn tool use driven through OpenCode; L3 is single-shot generation against the endpoint. Keeping them separate is deliberate: it lets “code generation is not the same as agentic tool use” be an empirical observation rather than an assumption.

Layer 1 runs on an ARM64-buildable 29-task subset of SWE-bench Verified, and we are explicit that this is *not* SWE-bench Verified performance and *not* a random sample. The official task images are x86-only; we build each task’s environment natively for aarch64 and keep only tasks whose environment builds and whose gold patch resolves on ARM64. This is an empirical buildability probe, not a curated difficulty sample. It biases the subset toward pure-Python repositories (flask, requests, pylint, pytest, sympy, sphinx, xarray, and similar) and excludes anything needing x86-only wheels or heavy C/Fortran toolchains—a coverage limitation we disclose. Because both models run the identical arm64-built tasks, any ARM-versus-x86 environment difference cancels in the relative comparison. The subset and its checksums are pinned in the provenance record.

Layer 2 has each model build the same realis-

tic full-stack application (“TaskFlow Local”) from an identical frozen specification, across a Node.js and a Python/FastAPI backend track. Scoring fires 29 equally weighted, fully automated HTTP assertions against a frozen API contract after booting the built app, reported as an acceptance-check fraction $k/29$. This is an API-acceptance signal, not a holistic app-quality score: it measures whether the app boots and satisfies the communicated contract, not user experience, design, or maintainability. We report a Bernoulli working-app rate (a per-run boot-plus-threshold pass) with a Wilson interval as the headline, never a single- N mean rubric score.

Layer 3 runs LiveCodeBench single-shot on a fixed pre-2024-06 window of 512 problems, scored as pass@1 with a Wilson 95% interval. A shared maximum generation budget of 8192 tokens is held identical across all three models as a fairness control—a decision that turns out to be consequential (Section 5.3).

Scoring and fairness. Scoring is deterministic given the retained artifacts, so runs can be rescored from committed outputs without re-serving. Same spec, prompt, hardware, time box, and harness for every run; all serving-setting differences that could not be equalized are disclosed; the maintainer’s tooling never edits model-generated application code, which lives in a separate repository.

5 Results

Figure 2 shows success by layer with Wilson intervals. On raw scores the ordering is $\text{gpt-oss} \geq \text{qwen} \geq \text{nemotron}$, but the L1 and L2 gaps are statistically weak and the L3 nemotron gap is budget-confounded. We report the numbers and then explain, layer by layer, why the durable finding is about how one measures rather than which model wins.

5.1 Layer 1: SWE-bench Verified ARM64 subset

Table 3 gives the Layer 1 rates. Two points about honest denominators. First, the raw rate charges operational failures—a watchdog kill and eight environment failures (clone/DNS/guard drops) that were never a measured model attempt—against the model. Separating them, which we track explicitly per run, raises nemotron’s rate from 20.7% to 30.0% and is the correct denominator for a *model* comparison. Second, pairwise McNemar tests on the shared 29 tasks find no significant separation among the three ($p = 0.125, 0.289, 1.0$). That is the appropriate result for a methodology framing and a fatal one for a ranking framing: at this sample size the layer cannot rank these models.

A second axis matters beyond “did it fix the bug.” SWE-bench also records whether a patch broke something that previously passed (PASS_TO_PASS). Among patches that applied and were evaluated, gpt-oss regresses at least one previously-passing test 33.3% of the time (8/24), qwen 31.8% (7/22), and nemotron 21.4% (3/14). The severity diverges far more than the rate: gpt-oss broke 1078 PASS_TO_PASS tests to nemotron’s 3, because a few catastrophic patches (for instance, breaking an import so a whole module’s suite fails) dominate the count. Figure 3 shows the trade-off: gpt-oss’s aggressive edits resolve the most and regress the most, while nemotron’s fewer, smaller edits rarely regress but rarely resolve. This characterizes failure quality among the non-resolving runs, not the winners, and the small denominators keep it directional.

5.2 Layer 2: app-build acceptance, and a harness bug that inverted the conclusion

The first Layer 2 sweep was produced under a harness-validity error we label C1. The frozen API contract that the prompt requires the model to satisfy had been dropped from the workspace by an app-repository tag-lineage break: four of the 29 checks were structurally unreachable, and every model was graded against a contract it could not see. After restoring the contract (app-repo tag `baseline-v7`) and re-running the vLLM models, the picture changed decisively.

Table 4 reports the corrected sweep. Under the bug, gpt-oss (node acceptance 0.252) and qwen (0.155) looked statistically indistinguishable—the working-app Wilson intervals overlapped, and the honest read was “too close to call.” Corrected, gpt-oss (0.724, 80% working apps) clearly and significantly beats qwen (0.178, 15%) on the node track, and the intervals separate. Crucially, the contract helped *only where a model could already build a booting app*: gpt-oss’s node score jumped $2.9\times$ (the four contract-only checks went from 0 to 48 passes), while qwen and both python tracks, which are floored for other reasons, barely moved. A single missing specification file had compressed a real $\sim 4\times$ capability gap into a fake statistical tie. Figure 4 shows the before/after.

The asterisk on nemotron marks an asymmetry we disclose rather than paper over. Nemotron was not re-run with the visible contract. Those four contract-only checks require a booted backend to be reachable, and nemotron almost never produced a working backend at Layer 2 (node acceptance mean 0.009, zero working apps). A visible contract is therefore mechanically a no-op for it, and a rerun would spend fragile TensorRT-LLM compute for no measurable effect. Its rows are retained and marked, and results are dual-reported as $k/29$ and a C1-rescored $k/25$ over the 25 reachable

Table 3: Layer 1, 29-task ARM64 subset, single-attempt pass@1. The “model-valid” rate excludes operational failures (a watchdog kill and eight clone/DNS/guard-drops that were never a measured attempt) from the denominator.

Model	Resolved (raw)	Model-valid rate	Notes
gpt-oss-120b	37.9% (11/29)	37.9% (11/29)	no infra failures
qwen3-coder-30b	24.1% (7/29)	24.1% (7/29)	no infra failures
nemotron-super	20.7% (6/29)	30.0% (6/20)	1 watchdog kill; 8 excluded

Table 4: Layer 2 after the C1 contract fix (contract-visible **baseline-v7**). Working-app rate is a per-run boot-plus-threshold pass; brackets are Wilson 95% intervals.

Model	node $k/29$	node working	python $k/29$
gpt-oss-120b	0.724	16/20 = 80%	0.069 (0/8)
qwen3-coder-30b	0.178	3/20 = 15%	0.004 (0/8)
nemotron-super*	0.009 (N=4)	0/4	0.000 (N=4)

checks (gpt-oss node 0.744, qwen 0.204). Two caveats survive the correction: the python track floors for both vLLM models, a genuine node-versus-python asymmetry rather than a contract artifact; and the layer remains an API-acceptance signal, not a full app-quality score.

5.3 Layer 3: LiveCodeBench, and a token budget that manufactured a last place

Table 5 gives raw Layer 3 pass@1, where nemotron finishes last. That last place is largely an artifact of the shared 8192-token budget, not of capability. Nemotron’s verbose reasoning exhausted the budget on 185 of 512 problems (143 empty outputs plus 42 with output but no extractable code), all scored as fails. The truncation is difficulty-correlated—no-code rate 14.3% on easy, 34.3% on medium, 71.5% on hard problems—so a conditional rate on only the answered subset must not be compared against the other models’ full-512 rates, a selection-bias trap we flag (finding C2). The defensible comparison is paired, on the 327 problems nemotron answered with code: there nemotron scores 96.0%, gpt-oss 95.4%, and qwen 82.3%. Figure 5 shows both halves.

The methodology finding is general: a single-shot, fixed-budget benchmark systematically penalizes verbose reasoning models, and the penalty concentrates on hard problems. We keep 8192 tokens as the primary, controlled number—holding the budget constant is the fairness control—and treat the truncation itself as a reportable result rather than tuning it away.

6 Measurement Pitfalls

Several pitfalls surfaced only because we retained raw artifacts and re-derived every table. They generalize beyond this box.

Unified memory is not VRAM. As noted, `nvidia-smi` reports [N/A] for device memory on GB10 by design; the footprint comes from `/proc/meminfo` and is whole-system, not dedicated GPU memory.

Tokens-per-joule is prompt-token-dominated. The agent loop re-sends context every turn, so prompt tokens are a median 98.7% of the total. A naive `total_tokens / energy` therefore measures re-submitted context, not decode work; we report energy per *generated* token alongside it.

Restart-truncated energy windows. Watchdog restarts reuse a run identifier and overwrite the metric window, so a crash-prone run’s energy can cover only its final segment. The gpt-oss Layer 3 energy cell is a known-broken capture (roughly three orders of magnitude low) and is marked invalid in the ledger rather than published—the pass@1 result is unaffected because scoring is independent of the metrics window.

Robust energy statistics. Arithmetic-mean Layer 1 energy is dominated by a single ~13 h nemotron operational runaway (~655 kJ). We report median and IQR with an operational-failure sensitivity, never a bare mean.

Telemetry asymmetry. vLLM exposes Prometheus histograms (gpt-oss ~26.9 tok/s, qwen ~19.7 tok/s at Layer 1); TensorRT-LLM exposes none, so nemotron has hardware energy but no decode-throughput or TTFT figures. Those cells are blank by design, not missing data, and we do not impute them.

7 Operational Findings

A few hard-won operational notes for anyone reproducing this. TensorRT-LLM 1.3.0rc9 on GB10 is fragile for non-NVIDIA MoE models—the FP4-native NVIDIA reasoning model serves, the others deadlock. A degraded-box state is real: after a 26-hour vLLM run the GPU was left in a state where TensorRT-LLM

Table 5: Layer 3 LiveCodeBench pass@1 (pre-2024-06 window, $n=512$). Nemotron’s raw rate is depressed by truncation, not capability (see text).

Model	pass@1	Wilson 95% CI	no-code / truncated
gpt-oss-120b	89.3%	[86.3, 91.7]	~1–2%
qwen3-coder-30b	68.2%	[64.0, 72.1]	~0%
nemotron-super	61.3%	[57.0, 65.4]	185/512

nemotron wedged after one generation (warmup 656 s); a reboot cleared it fully (warmup back to 8 s), so the box should be reset between heavy long-lived serves. A non-streaming request path wedged nemotron mid-generation, and switching the LiveCodeBench runner to stream (the path OpenCode already used) fixed it. Throughput is bandwidth-bound: nemotron’s single-stream Layer 3 would take ~130 h; running it 8-way parallel brings it to ~9.5 h with no change to per-problem pass@1 and no efficiency metrics lost, since TensorRT-LLM exposes none anyway. Finally, several silent-failure traps were guarded: reasoning models return `content=None` on truncation; a resume filter regenerated empty outputs indefinitely; and a watchdog nearly false-stalled on a frozen log mtime and now uses the TensorRT-LLM iteration counter as ground truth. Figures 6 and 7 summarize the efficiency reality, with the caveats above embedded.

8 Reproducibility

Every run keeps its summaries, manifests, and clock records under `results/raw/<run-id>/`; bulky per-run time series are regenerable and excluded from version control, while summaries and manifests are committed. Scoring is fully automated and mechanical—no human scorers, no LLM judge, at any layer—so the aggregation and robust-summary scripts regenerate their tables byte-identically from retained outputs. A single command (`make rebuild`) re-derives every summary table, per-run ledger, figure, and statistic from the raw runs without re-serving any model, and the statistics and robust-summary self-tests pass. Container image digests, upstream dataset snapshot revisions (SWE-bench Verified, LiveCodeBench), the exact HuggingFace revisions of all three model checkpoints, the Layer 1 subset checksums, and the frozen Layer 3 window are pinned in an immutable provenance record. Model-generated application code lives only in the companion app repository and is never edited by the maintainer’s tooling; the coding assistant used to build the harness and write this report is not part of what was measured.

9 Threats to Validity and Limitations

Single box, three configurations. Model identity is fully coupled to serving runtime, precision, and scaffold. Conclusions are restricted to these configurations on this machine, and serving incompatibility is a systems finding, not proof of intrinsic quality.

Small and unequal samples. L1 and L3 are single-attempt per cell; L2 is $N = 20/8$ (node/python) for the vLLM models against $N = 4$ for nemotron. All quality claims are descriptive, and the McNemar results show L1 cannot rank these models at this size.

Protocol changed mid-collection. The C1 contract fix and rubric-denominator changes mean some early runs are exploratory; confirmatory claims draw on the post-fix `baseline-v7` sweep.

L1 is a disclosed subset, not SWE-bench Verified and not a random sample; absolute SWE-bench numbers should not be extrapolated from it.

L3 is a fixed historical window where contamination is possible: equal chronological eligibility is not equal contamination, the set is not exposure-balanced, and it is not leaderboard-comparable.

Untested newer stacks. The nemotron/vLLM incompatibility is scoped to the pinned vLLM 0.15.1 (NGC 26.02-py3); a newer vLLM or the current NVIDIA-validated DGX Spark recipe was not tested and could change that specific result.

Isolation. Generated code is executed without OS-level sandboxing on this kernel (bubblewrap unavailable); publication-grade runs should use a no-network container or VM.

10 Conclusion

On a single DGX Spark the story that survives scrutiny is not “model X wins.” It is that at local-evaluation scale, serving feasibility and harness validity dominate model effects. No model in our set serves on both run-times on this box; a single missing specification file inverted the Layer 2 conclusion until it was found; and a fixed token budget manufactured a last-place finish at Layer 3 for a verbose reasoning model that matches the field on the problems it actually answers. The practitioner takeaways are concrete. Of the configurations

measured here, gpt-oss-120b on vLLM is the strongest local agentic setup (L1 37.9%, L2 node 80% working apps, L3 89.3%); nemotron serves only on TensorRT-LLM and needs a larger token budget to be judged fairly; and honest local benchmarking depends far more on a failure taxonomy, robust energy statistics, and byte-reproducible rescoring than on producing one more leaderboard number.

Data and Code Availability

The harness, frozen specification, raw run summaries, aggregation and statistics code, figures, and provenance pins are in the public repository at <https://github.com/mani-mal/spark-coder-bench>. Every table and figure in this paper maps to a committed source CSV and regenerates via `make rebuild`.

References

- [1] C. E. Jimenez, J. Yang, A. Wettig, S. Yao, K. Pei, O. Press, and K. Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? *ICLR*, 2024. arXiv:2310.06770.
- [2] N. Jain, K. Han, A. Gu, W.-D. Li, F. Yan, T. Zhang, S. Wang, A. Solar-Lezama, K. Sen, and I. Stoica. LiveCodeBench: Holistic and Contamination-Free Evaluation of Large Language Models for Code. arXiv:2403.07974, 2024.
- [3] W. Kwon, Z. Li, S. Zhuang, Y. Sheng, L. Zheng, C. H. Yu, J. E. Gonzalez, H. Zhang, and I. Stoica. Efficient Memory Management for Large Language Model Serving with PagedAttention. *SOSP*, 2023. arXiv:2309.06180.
- [4] W. Fedus, B. Zoph, and N. Shazeer. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity. *JMLR*, 2022. arXiv:2101.03961.
- [5] E. B. Wilson. Probable Inference, the Law of Succession, and Statistical Inference. *Journal of the American Statistical Association*, 22(158):209–212, 1927.
- [6] Q. McNemar. Note on the Sampling Error of the Difference Between Correlated Proportions or Percentages. *Psychometrika*, 12(2):153–157, 1947.
- [7] OpenCode. OpenCode: an open-source terminal-based coding agent. <https://opencode.ai>.
- [8] OpenAI. gpt-oss-120b model card. <https://huggingface.co/openai/gpt-oss-120b>, 2025.
- [9] Qwen Team. Qwen3-Coder-30B-A3B-Instruct model card. <https://huggingface.co/Qwen/Qwen3-Coder-30B-A3B-Instruct>, 2025.
- [10] NVIDIA. NVIDIA-Nemotron-3-Super-120B-A12B model card. <https://huggingface.co/nvidia/NVIDIA-Nemotron-3-Super-120B-A12B-NVFP4>, 2025.
- [11] Open Compute Project. OCP Microscaling Formats (MX) Specification (MXFP4; NVIDIA NVFP4). <https://www.opencompute.org/>, 2023.




Figure 1: End-to-end architecture of one benchmark run on a single DGX Spark. Sweep drivers read a single model registry and hand each arm its model plus serving profile; the selected model comes up behind an OpenAI-compatible endpoint (vLLM on :8000 or TensorRT-LLM on :8355) after a sanity gate and a per-serve reproducibility manifest; OpenCode drives the two agentic layers (L1, L2) while L3 calls the endpoint directly; synchronized collectors capture hardware energy/memory, vLLM Prometheus throughput, and agent accounting into per-run summaries that aggregate into the committed CSVs behind every table here.




Figure 2: Success by layer for the three configurations, each with a Wilson 95% interval: L1 resolved pass@1, L2 node working-app rate, L3 pass@1. The raw ordering is $\text{gpt-oss} \geq \text{qwen} \geq \text{nemotron}$, but the L1/L2 separations are weak and the L3 nemotron result is a token-budget artifact (Figure 5).




Figure 3: Layer 1 collateral damage, as a fraction of applied patches. gpt-oss both resolves the most and regresses the most; nemotron regresses least but resolves little. Aggressive editing buys resolutions at the cost of regressions.




Figure 4: The C1 ablation. Under the harness bug (contract invisible), gpt-oss (6/20) and qwen (4/20) working-app rates overlap. With the contract restored, gpt-oss jumps to 16/20 while qwen stays flat and the intervals separate. One missing spec file inverted the conclusion.




Figure 5: Left: nemotron's no-code (truncation) rate rises with difficulty (easy 14%, medium 34%, hard 72%). Right: on the 327 problems it answers with code, nemotron (96.0%) matches gpt-oss (95.4%) and beats qwen (82.3%). The last place in Table 5 is a fixed-budget artifact.




Figure 6: Median wall-clock per task, agentic layers. Nemotron is roughly 2× slower than the vLLM models at L1 (3.9 vs 1.4–1.9 min) and slowest at L2 (16.1 min); the median is shown because the L1 mean is skewed by one ~13h runaway.




Figure 7: Energy per *successful* task (self-hosted, so energy is the cost proxy). At L1, nemotron costs 143 kJ per solve to gpt-oss's 12.8 kJ (~11×); at L2, qwen costs 211 kJ per working app to gpt-oss's 32.8 kJ. Runtime-confounded (nemotron on TensorRT-LLM).

A Serving Commands and Configurations

Both runtimes are launched from a single sequential driver that writes a per-serve reproducibility manifest and passes a coherence sanity gate before any timed run. The vLLM invocation (Qwen shown; per-model values come from the `infra/models.json` registry, and `--max-num-seqs 1` is pinned across all models for clean per-task latency):

```
docker run -d --name vllm-server --gpus all --ipc=host \
--network=host --ulimit memlock=-1 --ulimit stack=67108864 \
-v $HF_HOME:/root/.cache/huggingface \
nvcr.io/nvidia/vllm:26.02-py3 \
vllm serve Qwen/Qwen3-Coder-30B-A3B-Instruct \
--host 0.0.0.0 --port 8000 --served-model-name qwen3-coder \
--max-num-seqs 1 --max-num-batched-tokens 16384 \
--dtype auto --trust-remote-code --seed 0 \
--enable-auto-tool-choice # + per-profile tool-call \
--parser
```

TensorRT-LLM serves Nemotron on port 8355 with an `extra_llm_api_options` YAML. The 1.3.0rc9 image is required: 1.2.1’s `nemotron_h` loader fell back to a regular init and double-allocated to ~115 GB, OOM-ing the host.

```
trtllm-serve nvidia/NVIDIA-Nemotron-3-Super-120B-A12B-NVFP4 \
--host 0.0.0.0 --port 8355 \
--max_batch_size 4 --tp_size 1 --ep_size 1 \
--max_num_tokens 8192 --max_seq_len 65536 \
--extra_llm_api_options /config/nemotron-super.yaml \
--reasoning_parser nano-v3 --tool_parser qwen3_coder
```

The YAML encodes the 128 GB unified-memory strategy for this hybrid Mamba/attention MoE (FP8 KV cache, float16 SSM cache with stochastic rounding, chunked prefill):

```
kv_cache_config:
  dtype: fp8
  enable_block_reuse: false
  free_gpu_memory_fraction: 0.9
  mamba_ssm_cache_dtype: float16
  mamba_ssm_stochastic_rounding: true
moe_config:
  backend: CUTLASS
cuda_graph_config:
  enable_padding: true
  max_batch_size: 8
enable_chunked_prefill: true
speculative_config:
  decoding_type: MTP
  num_nextn_predict_layers: 3
```

B The Nemotron vLLM Rejection

This is the headline serving-feasibility result of Section 4. The checkpoint’s `hf_quant_config.json` (produced by `modelopt 0.43.0.dev63`) declares a *mixed-precision* format—FP8 on the Mamba mixer projections, NVFP4 on the MoE experts, specified per layer:

```
{ "quantization": {
  "quant_algo": "MIXED_PRECISION",
  "kv_cache_quant_algo": "FP8",
  "quantized_layers": {
    "backbone.layers.0.mixer.in_proj": {"quant_algo": "FP8"},
    "backbone.layers.0.mixer.out_proj": {"quant_algo": "FP8"},
    "backbone.layers.1.mixer.experts.0.up_proj":
      {"quant_algo": "NVFP4", "group_size": 16},
    "...": "-every MoE expert listed, NVFP4 group_size 16"
  } } }
```

vLLM 0.15.1’s ModelOpt loader accepts only a single uniform algorithm and rejects the model during `create_engine_config(...)`—in ~6s, before any weight loads, and before any tool/reasoning-parser flag is reached:

```
pydantic_core._pydantic_core.ValidationError: 1 validation
error for VllmConfig
Value error, ModelOpt currently only supports:
['FP8', 'FP8_PER_CHANNEL_PER_TOKEN', 'FP8_PB_WO', 'NVFP4']
quantizations in vLLM. Please check the
'hf_quant_config.json' file for your model's quant config.
```

No profile or parser change fixes this: the rejected value is the checkpoint’s own format. A circulating workaround forces the top-level field to NVFP4; we did not use it, because the Mamba mixer projections are genuinely FP8 in the weights and forcing NVFP4 risks silently miscasting them. The same error against the same model and container is filed upstream at `vllm-project/vllm#37854` and `NVIDIA/dgx-spark-playbooks#77`.

C The Autonomous Tool-Use Gate

Before any timed agentic run, a model must pass a two-part gate; a model that fails is reported as a non-agentic baseline rather than scored as if it tried. (1) A direct request to the served endpoint with `tool_choice:"auto"` and a tool schema must return a non-null, well-formed `tool_calls` object—not prose, and not hallucinated tool-output tokens in the content. (2) In an OpenCode smoke test the model must actually create a file on disk. A model that only acts under `tool_choice:"required"` fails part (1). This is the gate DeepSeek-Coder-V2-Lite failed—it generated correct code but wrote nothing to disk while reporting success—and that `gpt-oss-120b` passed before replacing it.

D Layer 2 API Contract

Layer 2 scores 29 equally weighted, fully automated HTTP assertions after booting the built app. The checks and their categories:

```
boot: health_ok
auth: login_admin_ok, login_bad_creds_401,
      me_without_token_401, me_with_token_200,
```

```

login_member_ok
security: user_no_password_field,
protected_route_no_token_401
rbac: member_create_project_403
projects: admin_create_project, project_list_contains,
project_get_by_id, project_edit*, project_archive*
tasks: task_create, task_list, task_get,
task_update_status*, task_filter_status,
task_search, task_delete
comments: comment_add, comment_list
validation: validation_project_no_name_400,
validation_task_no_title_400,
validation_task_bad_priority_400
dashboard: dashboard_summary_keys*
frontend: frontend_build
reproducibility: lockfile_present

```

The four checks marked * are the C1 contract-only assertions of Section 5.2: they pin exact routes and key sets the requirements prose leaves unspecified, so they are reachable only when the frozen `api-contract.md` is in the workspace. Each check is a small deterministic assertion; the RBAC bypass check—which caught apps that let a non-admin member create a project—is representative:

```

# member (non-admin) must be forbidden from creating a project
r, _ = self.req("POST", C.API + "/projects",
                token=member_token, json={"name": "X"})
self.add("member_create_project_403", "rbac",
        r is not None and r.status_code == 403,
        r.status_code if r else "none")

```

E Layer 1 ARM64-Buildable Subset

The 29 tasks are every SWE-bench Verified instance whose environment builds natively on aarch64 and whose gold patch resolves there—an empirical buildability probe, not a difficulty sample. The distribution across the ten repositories that survived the probe:

Repository	N	Repository	N
astropy	4	pylint	3
django	3	pytest	4
matplotlib	2	requests	4
seaborn	2	sphinx	2
flask	1	sympy	4

```

astropy__astropy-12907  astropy__astropy-13033
astropy__astropy-13236  astropy__astropy-13398
django__django-10554    django__django-10880
django__django-10914    matplotlib__matplotlib-13989
matplotlib__matplotlib-14623  mwaskom__seaborn-3069
mwaskom__seaborn-3187    pallets__flask-5014
psf__requests-1142       psf__requests-1724
psf__requests-1766       psf__requests-1921
pylint-dev__pylint-4551  pylint-dev__pylint-4661
pylint-dev__pylint-4970  pytest-dev__pytest-10051
pytest-dev__pytest-10081  pytest-dev__pytest-10356
pytest-dev__pytest-5262  sphinx-doc__sphinx-10449
sphinx-doc__sphinx-10466  sympy__sympy-11618
sympy__sympy-12096       sympy__sympy-12419
sympy__sympy-12481

```

F Per-Model, Per-Layer Resource Metrics

Table 6 gives the consolidated hardware figures behind the efficiency discussion. Nemotron’s decode-throughput and TTFT cells are blank because TensorRT-LLM exposes no Prometheus histograms (Section 6), not because the data is missing. Layer 3 is omitted here: its duration is a whole-run figure rather than per-task, and the gpt-oss L3 metrics window failed to capture the generation (Section 6). Energy and wall-clock are per-task medians; memory is peak unified (system) RAM.

Table 6: Per-model, per-layer resource metrics (medians). Source: `results/summary/perf-resource-summary.csv`.

Layer	Model	N	min/task	kJ/task	GPU util %	GPU W	peak mem (MiB)	decode tok/s	TTFT p50 (s)
L1	gpt-oss-120b	29	1.9	3.6	90.7	30.1	113,059	26.2	0.47
L1	qwen3-coder-30b	29	1.4	2.6	89.3	29.2	110,455	20.1	0.50
L1	nemotron-super	21	3.9	8.4	91.8	38.0	108,517	—	—
L2	gpt-oss-120b	27	11.1	18.8	88.5	30.8	111,132	31.0	0.36
L2	qwen3-coder-30b	28	12.9	21.5	92.7	27.7	110,250	22.8	0.21
L2	nemotron-super	8	16.1	40.6	90.6	45.3	107,594	—	—